

Apache JClouds

Cloud interfaces, simplified

Hiranya Jayathilaka
Dept. of Computer Science, UCSB



Roadmap

- 💧 Cloud computing
- 💧 Challenges
- 💧 Apache JClouds
- 💧 Demo
- 💧 Pros and cons
- 💧 Summary and conclusion

The Cloud Revolution

- ◆ Cloud computing is revolutionizing the way software is developed and delivered.
 - ◆ Software-as-a-Service (SaaS)
 - ◆ Platform-as-a-Service (PaaS)
 - ◆ Infrastructure-as-a-Service (IaaS)

Many Benefits

- 💧 Cost effective
- 💧 Reduced maintenance overhead
- 💧 Easy and fast provisioning – Improved time-to-market
- 💧 Autoscaling and elasticity
- 💧 Fault tolerance

The Other Side of the Coin



Windows® Azure™



Challenges

- 💧 How do you select the cloud provider that's right for you?
- 💧 How do you port your application from one cloud provider to another?
- 💧 How to develop multi-cloud applications?

Enter, JClouds

- ◆ An open source library that facilitates developing applications for a wide range of cloud providers.
- ◆ Implement your application using JClouds, and run it on your favorite cloud without any code changes.
- ◆ Simple, intuitive APIs.
- ◆ Java and Clojure support.



History

- ◆ Started in March 2009 by Adrian Cole as an open source project.
- ◆ Initially based in London, developed mostly by European Java community.
- ◆ Contributed to Apache in April 2013.
 - ◆ <http://wiki.apache.org/incubator/jcloudsProposal>
- ◆ Already popular in the industry – Used by the likes of Adobe, CloudBees, RedHat, Twitter and Salesforce.
 - ◆ <http://jclouds.incubator.apache.org/documentation/reference/apps-that-use-jclouds/>

Getting Started

- ◆ No binary download ☺
- ◆ The documentation provides the necessary configurations for popular project management tools, so JClouds can be included in your project as a dependency.
 - ◆ Maven
 - ◆ ANT
 - ◆ Leiningen

JClouds APIs

- ◆ Compute Service API
 - ◆ For managing compute nodes (VMs) in the cloud
- ◆ Blobstore API
 - ◆ For storing data in the cloud

Compute Service API

- ◆ Key abstractions

- ◆ Hardware
- ◆ Operating system
- ◆ Template

- ◆ Supported providers

- ◆ AWS (EC2), Bluelock, CloudSigma, ElasticHosts, Go2Cloud, GoGrid, Green House Data, HP, Ninefold, OpenHosting, Rackspace, ServerLove, SkaliCloud, SoftLayer, Stratogen, TRMK, TryStack

Starting a VM from an Image

```
public static Set<? extends NodeMetadata> startVM(String region, String imageId, int count) {
    ComputeServiceContext context = ContextBuilder.newBuilder("ec2").
        credentials("my-access-key", "my-secret-key").
        modules(ImmutableSet.<~>of(new SshjSshClientModule())).
        buildView(ComputeServiceContext.class);

    ComputeService service = context.getComputeService();
    Template template = service.templateBuilder().
        imageId(region + "/" + imageId).
        smallest().build();

    try {
        return service.createNodesInGroup("jclouds", count, template);
    } catch (RunNodesException e) {
        handleException("Error starting the VMs", e);
    }
    return null;
}
```

If You Don't Have an Image?

```
public static Set<? extends NodeMetadata> startVM(String region, int count) {  
    ComputeServiceContext context = ContextBuilder.newBuilder("ec2").  
        credentials("my-access-key", "my-secret-key").  
        modules(ImmutableSet.<~>of(new SshjSshClientModule())).  
        buildView(ComputeServiceContext.class);  
  
    ComputeService service = context.getComputeService();  
    Template template = service.templateBuilder().  
        osFamily(OsFamily.UBUNTU).  
        osVersionMatches("12.04").  
        locationId(region).  
        smallest().  
        build();  
  
    try {  
        return service.createNodesInGroup("jclouds", count, template);  
    } catch (RunNodesException e) {  
        handleException("Error starting the VMs", e);  
    }  
    return null;  
}
```

More Control Over Templates

```
Template template = service.templateBuilder().  
    osFamily(OsFamily.UBUNTU).  
    osVersionMatches("12.04").  
    os64Bit(true).  
    minCores(2).  
    minRam(1024).  
    minDisk(2).  
    locationId(region).  
    build();
```

Access VM Metadata

```
public static void logNodeInfo(NodeMetadata node) {  
    log.debug("New node started successfully.");  
    log.debug("Node ID: " + node.getId());  
    log.debug("Node Image: " + node.getImageId());  
    log.debug("Node OS: " + node.getOperatingSystem().getFamily().value() +  
        "-" + node.getOperatingSystem().getVersion());  
    log.debug("Node Hostname: " + node.getHostname());  
  
    String s = "Private addresses: ";  
    for (String address : node.getPrivateAddresses()) {  
        s += address + "; ";  
    }  
    log.debug(s);  
  
    s = "Public addresses: ";  
    for (String address : node.getPublicAddresses()) {  
        s += address + "; ";  
    }  
    log.debug(s);  
}
```


SSH to Remote VM

```
ComputeServiceContext context = ContextBuilder.newBuilder("ec2").  
    credentials("my-access-key", "my-secret-key").  
    modules(ImmutableSet.of(new SshjSshClientModule())).  
    buildView(ComputeServiceContext.class);
```

```
NodeMetadata node = startNode(context, imageId);  
SshClient client = context.utils().sshForNode().apply(node);  
client.connect();  
ExecResponse response = client.exec("cat /etc/issue");  
if (response.getExitStatus() == 0) {  
    System.out.println(response.getOutput());  
} else {  
    System.err.println(response.getError());  
}  
client.disconnect();
```

Dealing with Package Managers

```
NodeMetadata node = startNode(context, imageId);
if (OperatingSystemPredicates.supportsApt().apply(node.getOperatingSystem())) {
    // Run apt-get commands
} else if (OperatingSystemPredicates.supportsYum().apply(node.getOperatingSystem())) {
    // Run yum commands
} else if (OperatingSystemPredicates.supportsZypper().apply(node.getOperatingSystem())) {
    // Run zypper commands
} else {
    // Default to something else - Or throw error
}
```

Managing Clusters

```
ComputeServiceContext context = ContextBuilder.newBuilder("ec2").
    credentials("my-access-key", "my-secret-key").
    modules(ImmutableSet.<~>of(new SshjSshClientModule())).
    buildView(ComputeServiceContext.class);
ComputeService service = context.getComputeService();
Template template = service.templateBuilder().
    imageId(region + "/" + imageId).
    minRam(4096).
    build();
service.createNodesInGroup("cassandra_cluster", 20, template);
service.runScriptOnNodesMatching(
    NodePredicates.runningInGroup("cassandra_cluster"),
    "/usr/local/cassandra/bin/run.sh start");

// Do work....

service.destroyNodesMatching(
    NodePredicates.runningInGroup("cassandra_cluster"));
```

Blobstore API

- ◆ Key abstractions

- ◆ Container

- ◆ Folder

- ◆ Blob

- ◆ Supported providers

- ◆ AWS (S3), CloudOne, HP, Azure, Ninefold, Rackspace, Synaptic

Downloading a Blob

```
public InputStream downloadBlob(String container, String fileName) {  
    BlobStoreContext context = ContextBuilder.newBuilder("aws-s3")  
        .credentials("my-access-key", "my-secret-key")  
        .buildView(BlobStoreContext.class);  
    BlobStore blobStore = context.getBlobStore();  
    Blob blob = blobStore.getBlob(container, fileName);  
    return blob.getPayload().getInput();  
}
```

Write Blob

```
public void writeBlob(String container, String fileName, String data) {  
    BlobStoreContext context = ContextBuilder.newBuilder("aws-s3")  
        .credentials("my-access-key", "my-secret-key")  
        .buildView(BlobStoreContext.class);  
    BlobStore blobStore = context.getBlobStore();  
    Blob blob = blobStore.blobBuilder(fileName).payload(data).build();  
    blobStore.putBlob(container, blob);  
}
```

Upload File

```
public void uploadFile(String container, File file) {  
    BlobStoreContext context = ContextBuilder.newBuilder("aws-s3")  
        .credentials("my-access-key", "my-secret-key")  
        .buildView(BlobStoreContext.class);  
    BlobStore blobStore = context.getBlobStore();  
    Blob blob = blobStore.blobBuilder(file.getName()).payload(file).build();  
    blobStore.putBlob(container, blob);  
}
```


Access Blob Metadata

```
MutableBlobMetadata metadata = blob.getMetadata();
System.out.println("Container: " + metadata.getContainer());
System.out.println("Name: " + metadata.getName());
System.out.println("ETag: " + metadata.getETag());
System.out.println("Date Created: " + metadata.getCreationDate());
System.out.println("Date Modified: " + metadata.getLastModified());
System.out.println("Provider: " + metadata.getProviderId());
System.out.println("URL: " + metadata.getPublicUri());

MutableContentMetadata contentMetadata = blob.getMetadata().getContentMetadata();
System.out.println("Size: " + contentMetadata.getContentLength());
System.out.println("MIME type: " + contentMetadata.getContentType());
System.out.println("Checksum: " + contentMetadata.getContentMD5());
System.out.println("Encoding: " + contentMetadata.getContentEncoding());
```

Logging Support

```
return ContextBuilder.newBuilder(properties.getProperty(CLOUD_PROVIDER)).  
    credentials(  
        properties.getProperty(CLOUD_ACCESS_KEY),  
        properties.getProperty(CLOUD_SECRET_KEY)).  
    modules(ImmutableSet.<~>of(  
        new Log4JLoggingModule(),  
        new SshjSshClientModule()))).  
    buildView(ComputeServiceContext.class);
```

Demonstration

◆ Scenario

- ◆ Start an Ubuntu VM in EC2
- ◆ Upload a Python script and a Java application to the VM (Java application also based on JClouds)
- ◆ Install Python and JRE on the VM using apt-get
- ◆ Run the Java application on the VM to download a data file from S3
- ◆ Run the Python script on the data file
- ◆ Upload the output to S3

Pros

- ◆ Powerful abstractions: Simple + High-level + Convention over configuration
- ◆ Location awareness baked into the API
- ◆ Easily manage clusters of nodes
- ◆ Excellent portability
- ◆ Unit testable
- ◆ Free and open source (ASL 2.0)
- ◆ “It just works” (YMMV)

Cons

- 💧 Still in the Apache Incubator
- 💧 Limited documentation
- 💧 Potential “Jar hell”

Summary

- 💧 Cloud computing brings a horde of benefits – But the diversity can be overwhelming.
- 💧 Several challenges with respect to evaluating cloud vendors and porting application across cloud platforms.
- 💧 Apache JClouds provides a simple and feature-rich approach for developing cloud applications that are easily portable across multiple vendors.

Thank You & Questions



References

- ◆ Project website: <http://jclouds.incubator.apache.org>
- ◆ Apache Incubator proposal:
<http://wiki.apache.org/incubator/jcloudsProposal>
- ◆ Getting started:
<http://jclouds.incubator.apache.org/documentation/gettingstarted/>
- ◆ Demo code: <https://github.com/hiranya911/jclouds-demo>